

In [1]: `import pandas as pd
import numpy as np`

In [3]: `from sklearn.datasets import fetch_california_housing
california_housing = fetch_california_housing()`

california_housing

```
{'data': array([[ 8.3252, 41., 6.98412698, ..., 2.55555556,
                 37.88, -122.23, ],
                [ 8.3014, 21., 6.23813708, ..., 2.10984183,
                 37.86, -122.22, ],
                [ 7.2574, 52., 8.28813559, ..., 2.80225989,
                 37.85, -122.24, ],
                ...,
                [ 1.7, 17., 5.20554273, ..., 2.3256351,
                 39.43, -121.22, ],
                [ 1.8672, 18., 5.32951289, ..., 2.12320917,
                 39.43, -121.32, ],
                [ 2.3886, 16., 5.25471698, ..., 2.61698113,
                 39.37, -121.24, ]]),
 'target': array([4.526, 3.585, 3.521, ..., 0.923, 0.847, 0.894]),
 'frame': None,
 'target_names': ['MedHouseVal'],
 'feature_names': ['MedInc',
                   'HouseAge',
                   'AveRooms',
                   'AveBedrms',
                   'Population',
                   'AveOccup',
                   'Latitude',
                   'Longitude']}
```

In [6]: `# creating a DataFrame`

```
df = pd.DataFrame(california_housing.data, columns=california_housing.feature_names)
df['MEDV'] = california_housing.target
df
```

Out[6]:

| | MedInc | HouseAge | AveRooms | AveBedrms | Population | AveOccup | Latitude | Longitude | MEDV |
|-------|--------|----------|----------|-----------|------------|----------|----------|-----------|-------|
| 0 | 8.3252 | 41.0 | 6.984127 | 1.023810 | 322.0 | 2.555556 | 37.88 | -122.23 | 4.526 |
| 1 | 8.3014 | 21.0 | 6.238137 | 0.971880 | 2401.0 | 2.109842 | 37.86 | -122.22 | 3.585 |
| 2 | 7.2574 | 52.0 | 8.288136 | 1.073446 | 496.0 | 2.802260 | 37.85 | -122.24 | 3.521 |
| 3 | 5.6431 | 52.0 | 5.817352 | 1.073059 | 558.0 | 2.547945 | 37.85 | -122.25 | 3.413 |
| 4 | 3.8462 | 52.0 | 6.281853 | 1.081081 | 565.0 | 2.181467 | 37.85 | -122.25 | 3.422 |
| ... | ... | ... | ... | ... | ... | ... | ... | ... | ... |
| 20635 | 1.5603 | 25.0 | 5.045455 | 1.133333 | 845.0 | 2.560606 | 39.48 | -121.09 | 0.781 |
| 20636 | 2.5568 | 18.0 | 6.114035 | 1.315789 | 356.0 | 3.122807 | 39.49 | -121.21 | 0.771 |
| 20637 | 1.7000 | 17.0 | 5.205543 | 1.120092 | 1007.0 | 2.325635 | 39.43 | -121.22 | 0.923 |
| 20638 | 1.8672 | 18.0 | 5.329513 | 1.171920 | 741.0 | 2.123209 | 39.43 | -121.32 | 0.847 |
| 20639 | 2.3886 | 16.0 | 5.254717 | 1.162264 | 1387.0 | 2.616981 | 39.37 | -121.24 | 0.894 |

20640 rows × 9 columns

In [7]: `df.info()`

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 20640 entries, 0 to 20639
Data columns (total 9 columns):
 #   Column      Non-Null Count  Dtype
---  -
 0   MedInc      20640 non-null  float64
 1   HouseAge    20640 non-null  float64
 2   AveRooms    20640 non-null  float64
 3   AveBedrms   20640 non-null  float64
 4   Population  20640 non-null  float64
 5   AveOccup    20640 non-null  float64
 6   Latitude    20640 non-null  float64
 7   Longitude   20640 non-null  float64
 8   MEDV        20640 non-null  float64
dtypes: float64(9)
memory usage: 1.4 MB
```

In [8]: `df.describe()`

Out[8]:

| | MedInc | HouseAge | AveRooms | AveBedrms | Population | AveOccup | Latitude | Longitude | MEDV |
|-------|--------------|--------------|--------------|--------------|--------------|--------------|--------------|--------------|--------------|
| count | 20640.000000 | 20640.000000 | 20640.000000 | 20640.000000 | 20640.000000 | 20640.000000 | 20640.000000 | 20640.000000 | 20640.000000 |
| mean | 3.870671 | 28.639486 | 5.429000 | 1.096675 | 1425.476744 | 3.070655 | 35.631861 | -119.569704 | 2.068558 |
| std | 1.899822 | 12.585558 | 2.474173 | 0.473911 | 1132.462122 | 10.386050 | 2.135952 | 2.003532 | 1.153956 |
| min | 0.499900 | 1.000000 | 0.846154 | 0.333333 | 3.000000 | 0.692308 | 32.540000 | -124.350000 | 0.149990 |
| 25% | 2.563400 | 18.000000 | 4.440716 | 1.006079 | 787.000000 | 2.429741 | 33.930000 | -121.800000 | 1.196000 |
| 50% | 3.534800 | 29.000000 | 5.229129 | 1.048780 | 1166.000000 | 2.818116 | 34.260000 | -118.490000 | 1.797000 |
| 75% | 4.743250 | 37.000000 | 6.052381 | 1.099526 | 1725.000000 | 3.282261 | 37.710000 | -118.010000 | 2.647250 |
| max | 15.000100 | 52.000000 | 141.909091 | 34.066667 | 35682.000000 | 1243.333333 | 41.950000 | -114.310000 | 5.000010 |

In [9]: `df.isnull().sum()`

Out[9]:

```
MedInc      0
HouseAge    0
AveRooms    0
AveBedrms   0
Population  0
AveOccup    0
Latitude    0
Longitude   0
MEDV        0
dtype: int64
```

In [11]: `# Seperating the features and target variable`

```
x = df.iloc[:, df.columns != "MEDV"]
y = df.iloc[:, df.columns == "MEDV"]
```

In [12]:

x

Out[12]:

| | MedInc | HouseAge | AveRooms | AveBedrms | Population | AveOccup | Latitude | Longitude |
|-------|--------|----------|----------|-----------|------------|----------|----------|-----------|
| 0 | 8.3252 | 41.0 | 6.984127 | 1.023810 | 322.0 | 2.555556 | 37.88 | -122.23 |
| 1 | 8.3014 | 21.0 | 6.238137 | 0.971880 | 2401.0 | 2.109842 | 37.86 | -122.22 |
| 2 | 7.2574 | 52.0 | 8.288136 | 1.073446 | 496.0 | 2.802260 | 37.85 | -122.24 |
| 3 | 5.6431 | 52.0 | 5.817352 | 1.073059 | 558.0 | 2.547945 | 37.85 | -122.25 |
| 4 | 3.8462 | 52.0 | 6.281853 | 1.081081 | 565.0 | 2.181467 | 37.85 | -122.25 |
| ... | ... | ... | ... | ... | ... | ... | ... | ... |
| 20635 | 1.5603 | 25.0 | 5.045455 | 1.133333 | 845.0 | 2.560606 | 39.48 | -121.09 |
| 20636 | 2.5568 | 18.0 | 6.114035 | 1.315789 | 356.0 | 3.122807 | 39.49 | -121.21 |
| 20637 | 1.7000 | 17.0 | 5.205543 | 1.120092 | 1007.0 | 2.325635 | 39.43 | -121.22 |
| 20638 | 1.8672 | 18.0 | 5.329513 | 1.171920 | 741.0 | 2.123209 | 39.43 | -121.32 |
| 20639 | 2.3886 | 16.0 | 5.254717 | 1.162264 | 1387.0 | 2.616981 | 39.37 | -121.24 |

In [13]:

y

Out[13]:

| | MEDV |
|-------|-------|
| 0 | 4.526 |
| 1 | 3.585 |
| 2 | 3.521 |
| 3 | 3.413 |
| 4 | 3.422 |
| ... | ... |
| 20635 | 0.781 |
| 20636 | 0.771 |
| 20637 | 0.923 |
| 20638 | 0.847 |
| 20639 | 0.894 |

In [14]:

x.shape

Out[14]:

(20640, 8)

In [15]:

y.shape

Out[15]:

(20640, 1)

In [16]: *# splitting the data into training(70%) and testing(30%) sets*

```
from sklearn.model_selection import train_test_split
x_train, x_test, y_train, y_test = train_test_split(x, y, test_size=0.3, random_state=123)
```

In [17]:

x_train

Out[17]:

| | MedInc | HouseAge | AveRooms | AveBedrms | Population | AveOccup | Latitude | Longitude |
|-------|--------|----------|----------|-----------|------------|----------|----------|-----------|
| 12364 | 3.9508 | 6.0 | 5.873188 | 1.111111 | 3010.0 | 3.635266 | 33.82 | -116.46 |
| 12271 | 3.5255 | 21.0 | 5.694581 | 1.049261 | 2174.0 | 2.677340 | 34.00 | -117.04 |
| 19605 | 1.9728 | 32.0 | 5.468208 | 1.144509 | 624.0 | 3.606936 | 37.55 | -121.03 |
| 10600 | 6.9133 | 8.0 | 5.976471 | 1.026471 | 862.0 | 2.535294 | 33.68 | -117.80 |
| 45 | 2.6768 | 52.0 | 4.335079 | 1.099476 | 718.0 | 1.879581 | 37.83 | -122.26 |
| ... | ... | ... | ... | ... | ... | ... | ... | ... |
| 7763 | 3.6389 | 36.0 | 5.584615 | 1.115385 | 490.0 | 3.769231 | 33.91 | -118.10 |
| 15377 | 4.5391 | 14.0 | 6.016688 | 1.017972 | 2436.0 | 3.127086 | 33.37 | -117.24 |
| 17730 | 5.6306 | 5.0 | 5.958393 | 1.031564 | 2435.0 | 3.493544 | 37.33 | -121.76 |
| 15725 | 3.8750 | 44.0 | 4.739264 | 1.024540 | 561.0 | 1.720859 | 37.78 | -122.44 |
| 19966 | 2.5156 | 20.0 | 5.491379 | 1.117816 | 1241.0 | 3.566092 | 36.21 | -119.08 |

In [18]:

x_test

Out[18]:

| | MedInc | HouseAge | AveRooms | AveBedrms | Population | AveOccup | Latitude | Longitude |
|-------|--------|----------|----------|-----------|------------|----------|----------|-----------|
| 19121 | 3.7917 | 40.0 | 4.959799 | 1.030151 | 1039.0 | 2.610553 | 38.24 | -122.64 |
| 20019 | 4.0217 | 9.0 | 5.804577 | 1.000000 | 1749.0 | 3.079225 | 36.09 | -119.05 |
| 15104 | 4.0882 | 12.0 | 5.360360 | 1.070571 | 3321.0 | 4.986486 | 32.85 | -116.98 |
| 3720 | 2.2377 | 27.0 | 3.376582 | 1.023207 | 3403.0 | 3.589662 | 34.20 | -118.42 |
| 8938 | 4.4211 | 41.0 | 5.656904 | 1.165272 | 1047.0 | 2.190377 | 34.01 | -118.47 |
| ... | ... | ... | ... | ... | ... | ... | ... | ... |
| 5462 | 2.5938 | 41.0 | 4.244444 | 1.148148 | 833.0 | 3.085185 | 33.99 | -118.47 |
| 1859 | 2.2750 | 21.0 | 5.457490 | 1.170040 | 1208.0 | 2.445344 | 41.95 | -124.14 |
| 10867 | 3.4835 | 17.0 | 3.825243 | 1.153099 | 2966.0 | 2.215086 | 33.70 | -117.88 |
| 4693 | 4.2500 | 52.0 | 5.736979 | 1.138021 | 899.0 | 2.341146 | 34.07 | -118.37 |
| 3521 | 4.5800 | 33.0 | 6.009862 | 1.039448 | 1578.0 | 3.112426 | 34.27 | -118.49 |

In [19]:

y_train

Out[19]:

| | MEDV |
|-------|-------|
| 12364 | 1.042 |
| 12271 | 1.321 |
| 19605 | 0.979 |
| 10600 | 2.741 |
| 45 | 1.823 |
| ... | ... |
| 7763 | 1.676 |
| 15377 | 1.809 |
| 17730 | 2.862 |
| 15725 | 4.125 |
| 19966 | 0.593 |

In [20]: y_test

Out[20]:

| | MEDV |
|-------|-------|
| 19121 | 1.516 |
| 20019 | 0.992 |
| 15104 | 1.345 |
| 3720 | 2.317 |
| 8938 | 4.629 |
| ... | ... |
| 5462 | 2.850 |
| 1859 | 1.224 |
| 10867 | 1.167 |
| 4693 | 4.869 |
| 3521 | 2.362 |

In [22]:

```
# creating the model for neural network architecture
# Sequential: it is linear stack of layers, created by adding layers one by one in sequence
# Dense: fully connected layer, adds dense layer to the model and takes 3 parameters(units, activation, name)

from keras.models import Sequential
from keras.layers import Dense

model = Sequential()
model
```

Out[22]: <Sequential name=sequential, built=False>

In [24]:

```
# this adds dense layers to the model

model.add(Dense(128, input_shape=(8, ), activation="relu", name="dense_1"))
model.add(Dense(64, activation="relu", name="dense_2"))
model.add(Dense(1, activation="relu", name="dense_output"))

D:\py\Lib\site-packages\keras\src\layers\core\dense.py:88: UserWarning: Do not pass an `input_shape`/`input_dim` argument to a layer. When using Sequential models, prefer using an `Input (shape)` object as the first layer in the model instead.
  super().__init__(activity_regularizer=activity_regularizer, **kwargs)
```

In [27]:

```
# compiling the model (setting the configuration for training it later on)
# takes 3 parameters(optimizer, loss, metrics)

model.compile(optimizer="adam", loss="mse", metrics=["mae"])
model.summary()
```

Model: "sequential"

| Layer (type) | Output Shape | Param # |
|----------------------|--------------|---------|
| dense_1 (Dense) | (None, 128) | 1,152 |
| dense_2 (Dense) | (None, 64) | 8,256 |
| dense_output (Dense) | (None, 1) | 65 |

Total params: 9,473 (37.00 KB)

Trainable params: 9,473 (37.00 KB)

Non-trainable params: 0 (0.00 B)

In [29]:

```
# training the neural network

history = model.fit(x_train, y_train, epochs=100, validation_split=0.05, verbose = 1)

Epoch 1/100
429/429 ----- 3s 2ms/step - loss: 514.7137 - mae: 4.6907 - val_loss: 5.4586 - val_mae: 2.0369
Epoch 2/100
429/429 ----- 1s 2ms/step - loss: 5.7319 - mae: 2.0945 - val_loss: 5.4586 - val_mae: 2.0369
Epoch 3/100
429/429 ----- 1s 2ms/step - loss: 5.8235 - mae: 2.1087 - val_loss: 5.4586 - val_mae: 2.0369
Epoch 4/100
429/429 ----- 1s 2ms/step - loss: 5.5772 - mae: 2.0657 - val_loss: 5.4586 - val_mae: 2.0369
Epoch 5/100
429/429 ----- 1s 2ms/step - loss: 5.7071 - mae: 2.0887 - val_loss: 5.4586 - val_mae: 2.0369
Epoch 6/100
429/429 ----- 1s 2ms/step - loss: 5.6863 - mae: 2.0820 - val_loss: 5.4586 - val_mae: 2.0369
Epoch 7/100
429/429 ----- 1s 2ms/step - loss: 5.7631 - mae: 2.0953 - val_loss: 5.4586 - val_mae: 2.0369
Epoch 8/100
429/429 ----- 1s 2ms/step - loss: 5.6668 - mae: 2.0806 - val_loss: 5.4586 - val_mae: 2.0369
Epoch 9/100
429/429 ----- 1s 2ms/step - loss: 5.6320 - mae: 2.0720 - val_loss: 5.4586 - val_mae: 2.0369
Epoch 10/100
429/429 ----- 1s 2ms/step - loss: 5.6320 - mae: 2.0720 - val_loss: 5.4586 - val_mae: 2.0369
```

In [30]:

```
mse_nn, mae_nn = model.evaluate(x_test, y_test)
print("mse_nn: ", mse_nn)
print("mae_nn: ", mae_nn)

194/194 ----- 0s 1ms/step - loss: 5.6806 - mae: 2.0799
mse_nn: 5.541019916534424
mae_nn: 2.0531160831451416
```

In []: